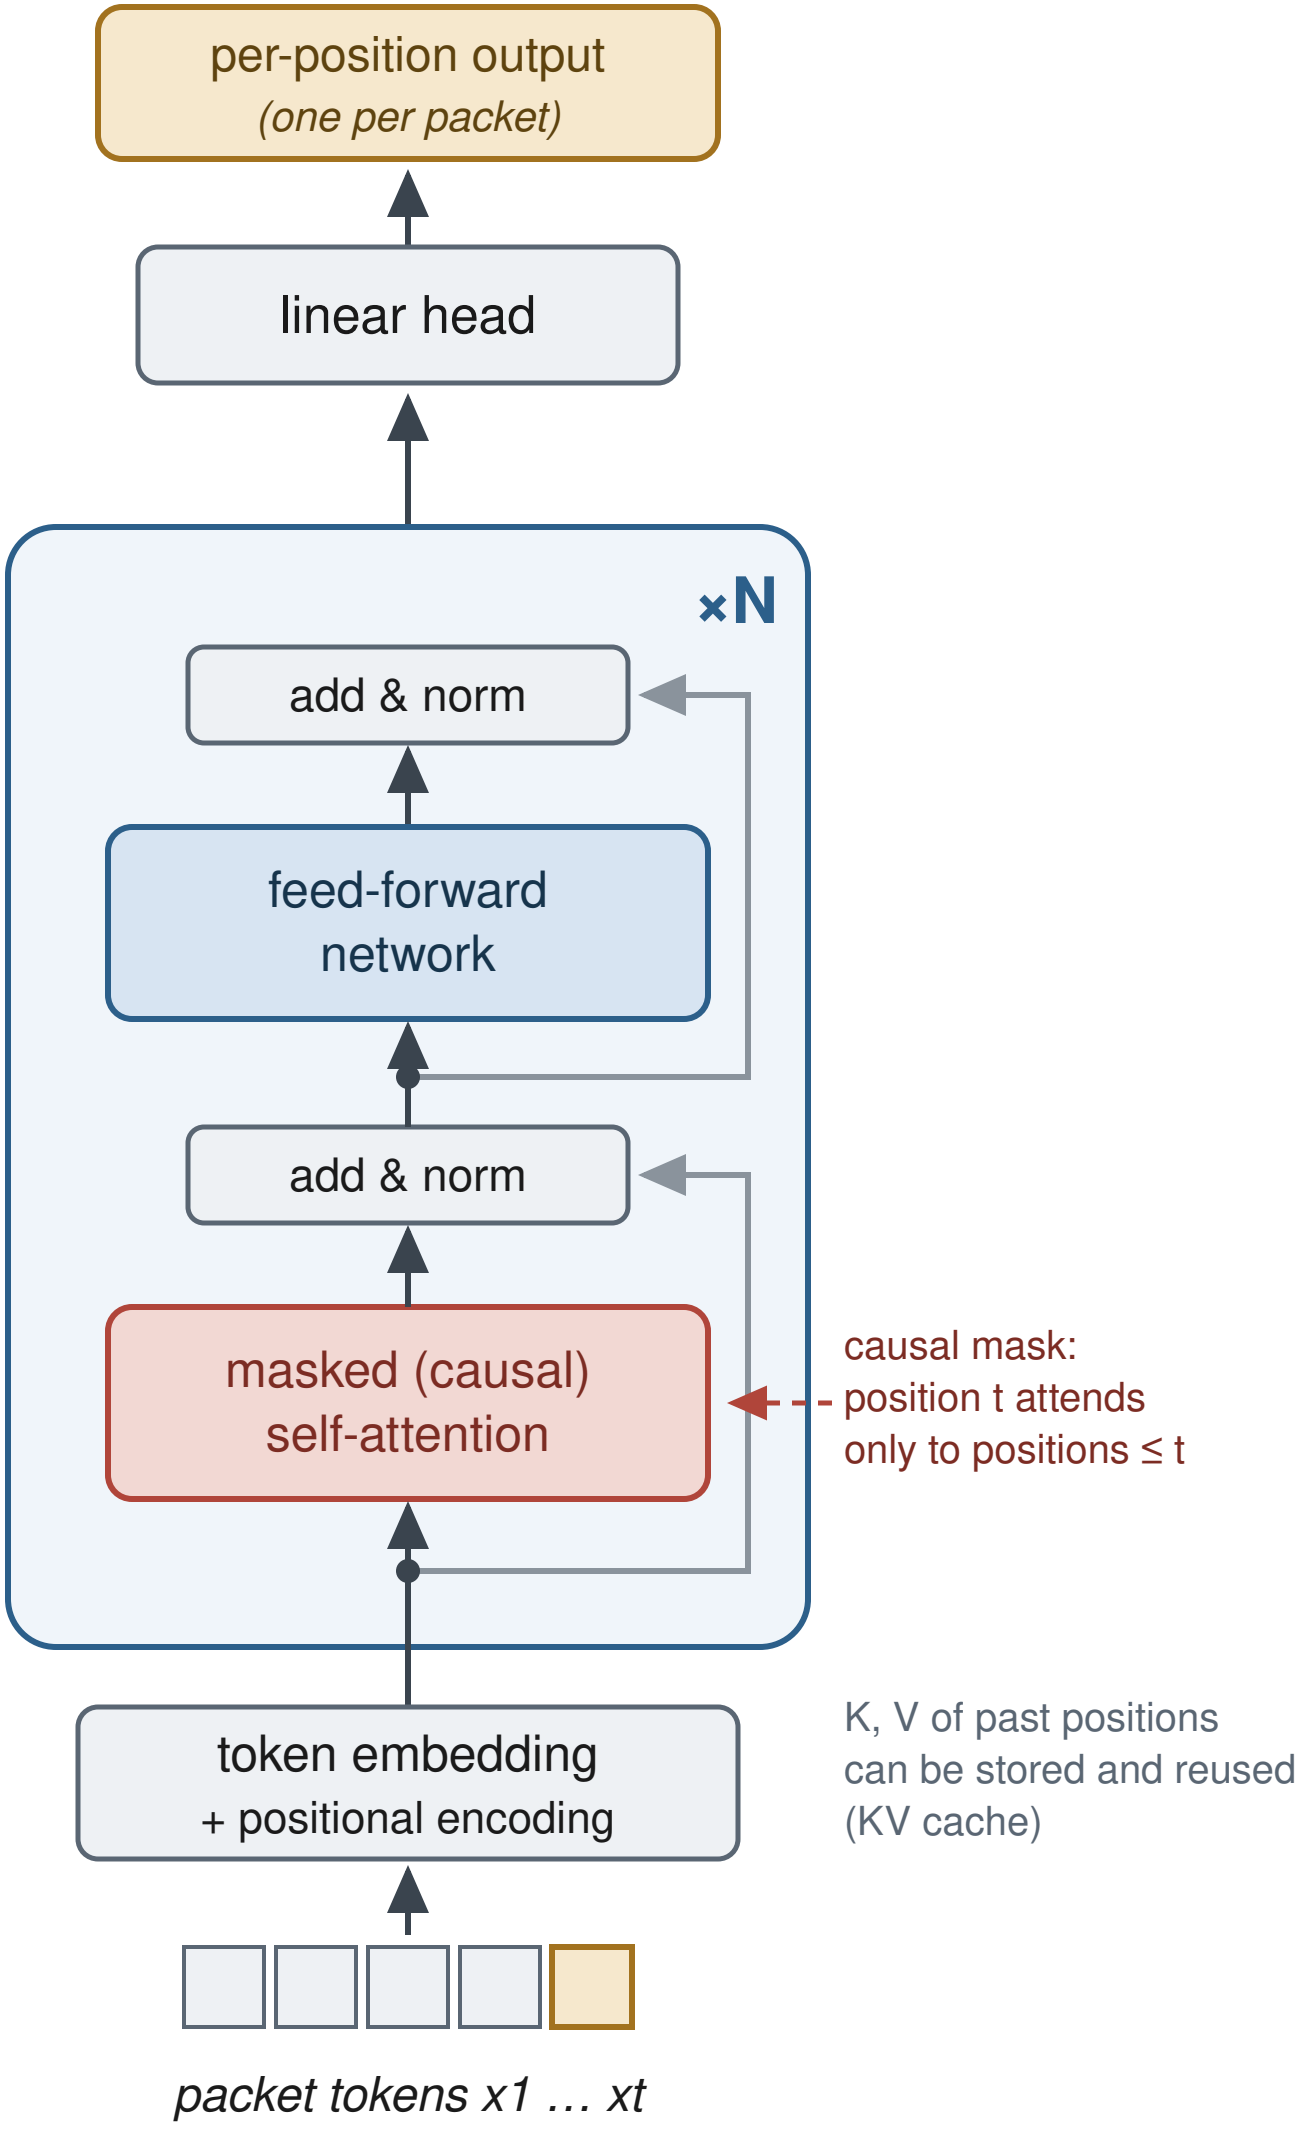
per-position output
(one per packet)
linear head
×N
add & norm
feed-forward network
add & norm
masked (causal) self-attention
causal mask:
position t attends
only to positions ≤ t
token embedding
+ positional encoding
K, V of past positions
can be stored and reused
(KV cache)
packet tokens x1 … xt